

Week 3 - System Calls

Welcome to week 2! This week we'll be talking about system calls, also known as syscalls. This is the main interface your programs use to interface with your OS, and they can provide a lot of very valuable information about what your programs are doing. Without them, they're just a black box that doesn't provide any input or output to the outside world. Here we'll go over why they exist, how they work, and some examples. Let's get started!

The process abstraction

As we talked about last week, the OS is in charge of managing access to hardware, and perform hardware-level operations. These devices include things like the hard drive, network card, the USB ports (which could have things like a mouse, keyboard, camera, etc. connected to it), audio input and output, GPU's, among many, many others. However, the OS has to ensure that your programs can't access these things directly, and so it provides a very basic abstraction: the process.

The process abstraction is a way your OS lets you encapsulate your programs. It gives you a work area where, as far as your program can tell, you have access to the whole of contiguous memory, every possible address out there, and it makes it look like it's all owned only by your process. It makes it seem like you have full control over what's being processed on this part of your system. These programs all run in what we like to call "user mode". This means that a lot of features and instructions you'd be able to use on your CPU from the kernel are completely inaccessible here (things like direct hardware access, modifying memory, setting up and modifying interrupts, etc.).

You might wonder, why do we want this? Well, there are a couple of advantages this gives us:

1. **Isolation:** Every process on your system cannot affect any other process (except via well-defined channels, more on that later).
2. **Reality-abstraction:** Real-life constraints are messy. Physical memory runs out, changes from machine to machine, processes get pre-empted. The memory and CPU abstractions allow us to ignore these limitations, pretending we have infinite memory and access to CPU when we write the program, and we can let the OS figure out how to provide that for us.
3. **No direct HW access:** This is the most important one. Hardware access needs to be coordinated between the different processes. For example, you need to pick where to point an HDD's spindle to read a sector. You could leave this coordination up to each process through some sort of library and let the user programs send commands straight to the HDD. However, this would let a program affect and interrupt other programs, so instead the OS just takes away this ability together.

Note on CPU rings

One question that often comes up, is how does the OS block access to certain CPU instructions? After all, couldn't you just put the right thing in your program and it would access the hard drive without your OS knowing? Indeed, you can do a lot of things without your OS ever finding out. However, most modern OS' make use of a feature provided by most modern CPU's: protection rings.

You can read more about them here, but in summary, most CPU's provide multiple "modes" you can run in, with different restrictions in each mode. Each mode lets you run different instructions, and only the more privileged modes will let you access hardware, modify memory,

etc. On boot, your OS will start in the highest privilege, and set up everything so that, whenever it launches new processes, they execute in the lowest privilege possible. As such, the kernel itself always runs code in the highest privilege.

Notice this is different from the concept of “running as root”. Even root programs run in that lowest privilege ring.

Interfacing with the kernel

One thing you’ll notice in the process abstraction is that, by itself, it shouldn’t let you get anything done. You can’t access your hardware, you can’t read and write to disk (or files in general for that matter), and you can’t even write to the terminal. However, as we saw last week, you can run plenty of programs in Linux that do all these things even though all those programs are running as these sandboxed processes. How come?

Simple: system calls. System calls (or syscalls) are a mechanism that Linux (and most other operating systems) provide for interfacing with the kernel. They’re special functions that will let you pass information down to the kernel and ask it to execute a well-defined operation. They include operations like `open()` for opening a file, `read()` for reading a specified number of bytes out of an open file, `fork()` for creating a copy of the current process into a new one, `stat()` to get metadata on a file, `kill()` for sending signals to another process (some of which can kill them), etc. There’s a huge set of system calls, but for this module, we’ll focus on understanding how they work at a high level, and we’ll use some of the basic file manipulation syscalls.

Continuing...

We’ll go into more detail about how syscalls work and some implementation details in the exercises. Each exercise will have its image, much like in the last week’s problems. Go through each exercise PDF in order (titled `ex_1.pdf`, `ex_2.pdf`, and so on), and then take a look at the `problems.pdf` for open-ended problems to apply what you learned.

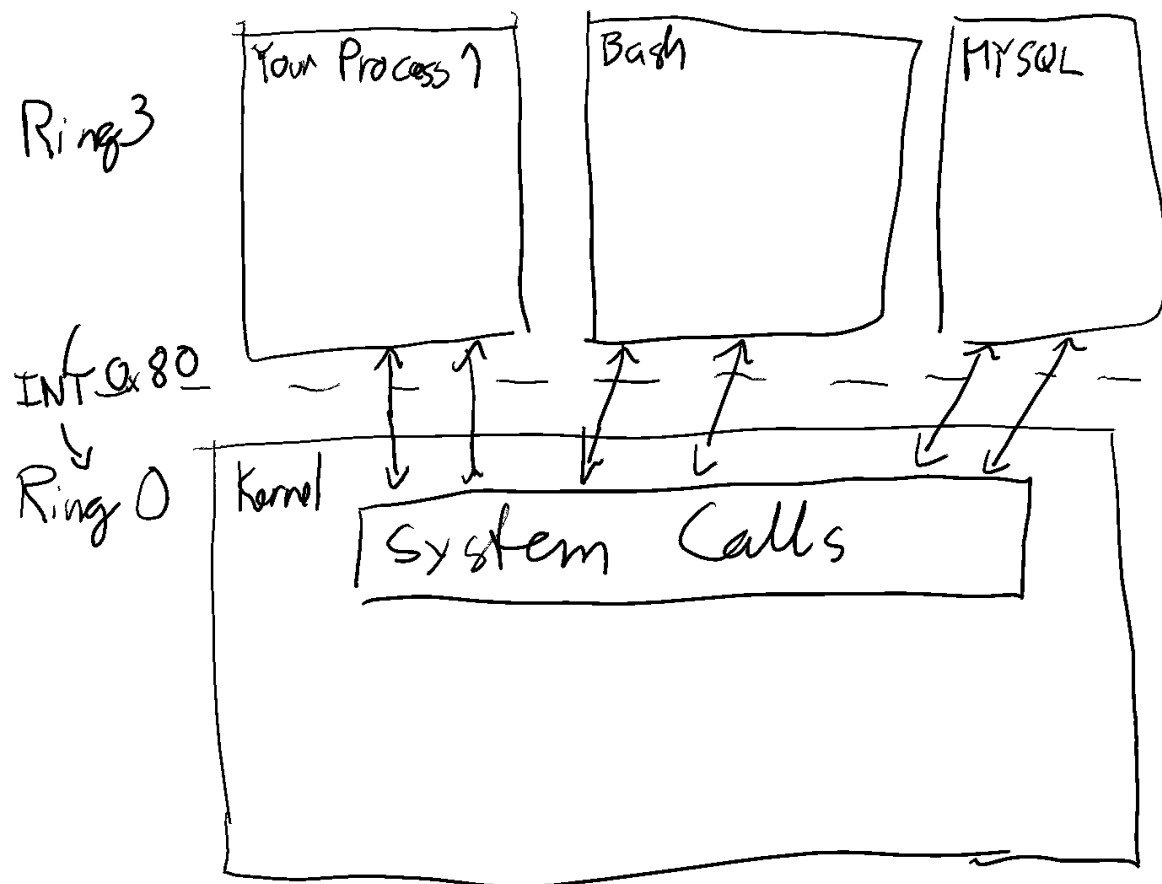


Figure 1: A brief, hand-drawn summary of syscalls